

# AI Chatbot Testing Strategy

## Context:

The product is an AI-powered chatbot that generates non-deterministic responses. The same input may produce different, yet valid, outputs. Therefore, traditional exact-match validation is not sufficient. The testing approach must focus on semantic correctness, reliability, performance, and risk mitigation.

---

## 1. Validation Strategy for Non-Deterministic Outputs

Since exact output matching is not possible, validation must be outcome-based rather than string-based.

### A. Intent & Semantic Validation

Instead of comparing exact text, validate:

- Whether the response correctly addresses the user's intent
- Whether key entities or expected concepts are present
- Whether the response aligns with the prompt's objective

Example:

Prompt: "Summarize the following paragraph."

Validation checks:

- Output length < input length
- No new facts introduced
- Main keywords/entities preserved
- Summary captures core meaning

This can be implemented using:

- Keyword validation
  - Rule-based assertions
  - Semantic similarity scoring (embeddings comparison)
  - Human-reviewed golden dataset for baseline validation
- 

## B. Rule-Based Guardrails

Define expected behavior rules per feature:

Examples:

- If user asks for a list → response must contain bullet points or numbered items
- If user asks for calculation → numeric output must be mathematically correct
- If user asks about restricted content → response must comply with safety policy

These rules allow partial deterministic validation.

---

## C. Multi-Run Consistency Testing

Because responses vary, the same prompt should be executed multiple times (e.g., 5–10 runs).

Measure:

- Intent consistency
- Semantic similarity range
- Variance score

This helps detect unstable model behavior.

---

# 2. Key Metrics to Test

## 1. Confidence Score

If the model provides a confidence score:

- Validate that it meets minimum threshold (e.g.,  $> 0.7$ )
  - Flag low-confidence responses
  - Compare confidence drift across versions
- 

## 2. Latency

AI systems are sensitive to response time.

Test:

- Average response time
- Timeout handling
- Performance under concurrent load

Set acceptable SLA (e.g.,  $< 3$  seconds for standard queries).

---

## 3. Token Usage

Since token usage affects cost and performance:

Validate:

- Input tokens
- Output tokens
- Total token consumption per request
- Unexpected token spikes

Monitor token efficiency across releases.

## 4. Hallucination Detection

Hallucination is a critical AI risk.

Approach:

- Use factual benchmark dataset
- Cross-verify output with trusted source
- Flag unsupported claims
- Check for fabricated citations or data

Example:

If chatbot answers factual questions, compare output against known verified dataset.

---

## 5. Safety & Compliance

Test for:

- Harmful content generation
- Bias
- Policy violations
- Prompt injection vulnerabilities

Include adversarial testing scenarios.

---

## 3. Regression Testing Strategy for AI Features

AI regression differs from traditional regression.

## A. Golden Dataset Approach

Create a fixed dataset of:

- Critical prompts
- Business-sensitive queries
- Edge cases
- Safety scenarios

For each release:

- Run all prompts
- Compare:
  - Semantic similarity
  - Accuracy classification
  - Latency
  - Token usage

Track deviations.

---

## B. Version Comparison Testing

When model or prompt logic changes:

Compare:

- Previous model vs new model
- Intent accuracy rate
- Hallucination rate

- Latency difference
- Cost impact (tokens)

Create a regression dashboard for trend monitoring.

---

## C. Drift Detection

Monitor:

- Changes in output distribution
- Drop in similarity score
- Increase in hallucination frequency

Automated alerts can detect degradation early.

---

## D. Human-in-the-Loop Validation

For critical workflows:

- Periodic manual review of sampled responses
- Feedback loop into evaluation framework
- Continuous improvement process

# 4. Risk-Based Prioritization

Focus testing effort on:

- High-impact business flows
- Legal or compliance-sensitive prompts
- Financial or medical queries

- Customer-facing production flows